

Programming Assignment 3

Due Wednesday, October 28 at 6 pm

Dijkstra's Algorithm

Edsger Dijkstra was a Dutch computer scientist. He is unquestionably one of the founding fathers of computer science. In addition to devising Dijkstra's algorithm, he was one of the originators of *structured programming* — a precursor to modern object-oriented programming. C, for example, is a structured programming language. He also made *many* contributions to the solution of process-coordination problems in parallel computing: we'll learn about several later on in the course. His Wikipedia page lists more than fifty fundamental contributions to computer science.

Dijkstra's algorithm solves “the single-source shortest path problem” in a weighted, directed graph. The input consists of a weighted, directed graph and a specified vertex in the graph. A *directed graph* or *digraph* is a collection of *vertices* and *directed edges*, which join one vertex to another. A *weighted digraph* has numerical weights assigned to the vertices and/or the edges.

The input graph for Dijkstra's algorithm has nonnegative weights assigned to the edges, and the algorithm finds a shortest path from the specified vertex, the *source* vertex, to every other vertex in the graph. The algorithm uses four main data structures:

1. The adjacency matrix **mat** of the graph: **mat**[**v**][**w**] is the cost or weight of the edge from vertex **v** to vertex **w**. If there is no edge joining the vertices, the cost or weight is ∞ .
2. A list **dist**: **dist**[**v**] is the current estimate of the cost of the least cost path from the source to **v**.
3. A list **pred**: **pred**[**v**] is the “predecessor” of vertex **v** on the current path 0- \rightarrow **v**.
4. A set **known** consisting of those vertices for which the shortest path from the source to the vertex is known.

Our vertices will be $0, 1, 2, \dots, n - 1$, and vertex 0 is the source.

The algorithm is divided into phases: initialization and path-finding. During initialization, **dist**[**v**] is set to **mat**[0][**v**] for each vertex **v**. So our initial estimate of the length of the shortest path 0- \rightarrow **v** is just the length of the edge from 0 to **v**. Of course, if there is no edge 0- \rightarrow **v**, then **dist**[**v**] will be ∞ . Initially, then 0 is the “predecessor” on the path 0- \rightarrow **v**. So **pred**[**v**] should be set to 0 for each vertex **v**. We'll require that the distance from any vertex to itself is 0. So **mat**[**v**][**v**] is 0 for every vertex **v**. and after the initialization, 0 will

have `dist[0] = 0` and `pred[0] = 0`. Since edges have nonnegative weights, we know the shortest path to 0. So `known = {0}`.

After initialization, we iterate $n - 1$ times, once for each remaining vertex. In each iteration, we first find a vertex `u` whose distance to 0 is as small as possible among the vertices not already in the set `known`. This vertex `u` is added to `known`, and the `dist` and `pred` entries are updated for the vertices still not in `known`: if there is a shorter path to `v` obtained by going through `u`, then `dist` and `pred` are updated accordingly.

This gives the following algorithm:

Dijkstra's Algorithm

Input: `n`: the number of vertices in the digraph

`mat`: the adjacency matrix of the digraph:

`mat[v][w]` is the length of the edge from `v` to `w` or infinity if there is no edge from `v` to `w`.

```
for (v = 0; v < n; v++) {
    dist[v] = mat[0][v]; /* Length of edge 0->v or infinity */
    pred[v] = 0;         /* 0 precedes v on the current shortest path */
    known[v] = 0;        /* We don't know the shortest path to v */
}
known[0] = 1;           /* We do know the shortest path to 0 */

for (int i = 0; i < n; i++) {
    /* Among vertices not in known, find vertex */
    *    u with smallest dist val                */
    u = Find_min_dist(dist, known, n);

    /* Add u to known */
    known = known + {u}.

    /* Among vertices not in known, update Dist[v]      */
    /* and Pred[v] if there's a shorter path 0->u->v */
    for (v = 1; v < n; v++)
        if (v is not in known) {
            new_dist = dist[u] + mat[u][v];
            if (new_dist < dist[v]) {
                dist[v] = new_dist;
                pred[v] = u;
            }
        }
    } /* if v */

} /* for i */
```

Output:

Length of shortest path 0→v for each vertex v

Shortest path 0→v for each vertex v

Parallelizing Dijkstra's Algorithm

We can divide the work among several processes by assigning the responsibility for different vertices to different processes. If p , the number of processes evenly divides, n , the number of vertices, then we can assign responsibility for the vertices using a *block* partition:

```
Process 0: Vertices 0, 1, ..., n/p - 1
Process 1: Vertices n/p, n/p+1, ..., 2n/p - 1
...
Process p-1: Vertices (p-1)n/p, (p-1)n/p + 1, ... , pn/p-1 = n-1
```

For example, if $p = 3$, and $n = 6$, then the vertices should be partitioned as follows:

```
Process 0: Vertices 0, 1
Process 1: Vertices 2, 3
Process 2: Vertices 4, 5
```

Most of the computational work in the program is in the code that finds **u**, the vertex with minimum **dist** value, and the code that updates the **dist** and **pred** values after a new vertex has been added to **known**. In order to find **u** in the serial algorithm, we can do something like this:

```
u = -1;
min_dist = INFINITY;
for (v = 0; v < n; v++)
    if (v is not in known)
        if (dist[v] < min_dist) {
            u = v;
            min_dist = dist[v];
        }
```

So if **dist** and **known** are divided among the processes, this can be parallelized with essentially the same code:

```
loc_u = -1;
loc_min_dist = INFINITY;
for (loc_v = 0; loc_v < loc_n; loc_v++)
    if (loc_v is not in loc_known)
        if (loc_dist[loc_v] < loc_min_dist) {
```

```

        loc_u = loc_v;
        loc_min_dist = loc_dist[v];
    }

```

Here, `loc_n = n/p`. Of course, now each process will have a candidate for `u`. So to get the “global” minimum vertex, we need to take a global minimum. This can be done fairly efficiently with tree-structured communication. However, it’s easier to use the MPI function `MPI_Allreduce`. To do this each process should declare two two-element arrays of `int`:

```
int my_min[2], glbl_min[2];
```

The values in `my_min` should be

```
my_min[0] = loc_dist[loc_u];
my_min[1] = Global_vertex(loc_u);
```

So the first element is the distance from 0 to `loc_u`. To understand the second, note that `loc_u` is a *local* vertex number: it ranges from 0 to `loc_n = n/p`. It needs to be a *global* vertex number. We can compute this by simply computing

```
global_u = loc_u + my_rank*loc_n;
```

Now we can call `MPI_Allreduce`:

```
MPI_Allreduce(my_min, glbl_min, 1, MPI_2INT,
              MPI_MINLOC, comm);
```

This will return the global min dist in `glbl_min[0]` and the global vertex number of the vertex with this min dist in `glbl_min[1]`. So

```
u = glbl_min[1].
```

A few comments on the call to `MPI_Allreduce` are in order: First we pass in both the minimum distance on each process together with the vertex where that minimum distance occurs. Our call to `MPI_Allreduce` will find the minimum of all the distances, and store that in `glbl_min[0]`. It will also find the vertex where the global minimum occurred and store that in `glbl_min[1]`. At this point we’re really interested in `u`, the vertex where the global min occurs, but the distance is needed to find `u`. So we pass in a two-element array of ints. Hence, we use the special MPI datatype `MPI_2INT`, which can be used in communications for arrays storing two ints.

For the second part of the main loop, we need to look at all the vertices assigned to the current process that are not yet elements of `known`. For each such vertex `loc_v`, we need to compare its current distance from 0 to the distance obtained by looking at the path to `loc_v` that goes through `u`, the vertex just added to `known`. If each process has all the *columns* of `mat` that correspond its vertices, then this “submatrix” of columns can be used to calculate the possibly new distance:

```
new_dist = dist_u + loc_mat[u][loc_v];
```

Then if `new_dist` is smaller, we can update the current distance and the current predecessor:

```
if (new_dist < loc_dist[loc_v]) {
    loc_dist[loc_v] = new_dist;
    loc_pred[loc_v] = u;
}
```

Program 3

For programming assignment 3, you should write an MPI program that parallelizes Dijkstra’s algorithm using the approach discussed in the previous section. Input to the program will consist of the following.

1. The number n of vertices in the digraph. You can assume that n is evenly divisible by p , the number of processes.
2. The “adjacency matrix” `mat` of the digraph. This is a matrix with n rows and n columns. The entry in row v and column w , `mat[v][w]`, is the distance *from* vertex v *to* vertex w . All the entries will be nonnegative ints, and the entries on the “diagonal” (row subscript = column subscript) will be 0.

The output of the program should consist of

1. A list of the lengths of the shortest paths from vertex 0 to each vertex v , and
2. The vertices on the shortest $0 \rightarrow v$, for each vertex v .

Implementation Details

Working with arrays in distributed memory programs can be confusing, since subscripts may be different from the corresponding serial program. For example, suppose we have the following 4×4 matrix `mat`

```
0 3 1 2
5 0 9 1
6 1 0 4
4 2 8 3
```

in a serial program, but the matrix is divided among two processes in a parallel program:

```
Proc 0:  0 3      Proc 1:  1 2
          5 0          9 1
          6 1          0 4
          4 2          8 3
```

Then in the serial program, the element in row 2, column 3 is accessed with `mat[2][3]`. But if the process submatrices are called `loc_mat`, then this element is stored in row 2, column 1 on process 1: `loc_mat[2][1]`. In this program row subscripts will be the same in both the parallel and serial programs, but column subscripts will be different.

To convert a “local” column or vertex to a “global” column or vertex, you can use the formula:

```
global vertex = local vertex + my_rank*loc_n
```

To go the other way

```
local vertex = global vertex % loc_n
```

To determine whether a process has been assigned a particular global vertex, you can check whether

```
(global vertex)/loc_n == my_rank
```

Although C does have two-dimensional arrays, there are serious limitations on how they can be used. So any matrices used in your program (i.e., the input matrix `mat` and the block-column submatrices `loc_mat`) should be stored as one-dimensional arrays, and elements should be accessed using arithmetic:

The entry in row `v` and column `w` of `mat` is `mat[v*n + w]`.

The entry in row `v` and column `loc_w` of `loc_mat` is

```
loc_mat[v*loc_n + loc_w].
```

Storage for these matrices can be allocated with `malloc`:

```
loc_mat = malloc(n*loc_n*sizeof(int));
```

This will allocate storage for `n*loc_n` ints, and elements can be accessed using ordinary array subscripts. A similar approach can be used to allocate storage for `dist`, `loc_dist`, etc. Of course, storage allocated with `malloc` should be freed when it’s no longer needed:

```
free(loc_mat);
```

will free all the storage allocated with the call to `malloc`.

The global constant `INFINITY` will be input as 1000000. You can use some other value in your program, but your program must recognize 1000000 as `INFINITY` on input.

There some I/O functions on the class website in the file `mpi_io.c`. The most important of these is `Read_Matrix`. It uses a special MPI Datatype to distribute the matrix columns among the processes. You can treat this function and its supporting function `Build_blk_col_type` as a black box: you don’t need to learn how they work.

The distributed arrays `dist` and `pred` can be collected onto process 0 for processing and/or output using the MPI collective communication function `MPI_Gather`. For example, if `dist` consists of n ints with a block of $n/p = \text{loc_n}$ ints assigned to each process in `loc_dist`, then `dist` can be collected onto process 0 with the call

```
MPI_Gather(loc_dist, loc_n, MPI_INT, dist, loc_n, MPI_INT,
          0, comm);
```

Note that the arg following `dist, loc_n`, is the number of elements received from *each* process. It's not the total number of elements received.

Development

You can compile and test your program using any system with an MPI implementation. However, before your final submission, you should check your program on the penguin cluster, since this is where we will test your program.

Using conventional debuggers with distributed memory programs can be quite difficult: it will probably be easier to debug your code using `printf/fflush`. Remember that *any* process can print to `stdout`. So any process can print output. However, the order in which output is actually printed may be fairly random. In fact, the output from one process can be interrupted by another process. The chances that this will happen are somewhat less, if you put a block of your output into a string, and then just print the string. For an example of this, see the function `Print_local_matrix` in the file `mpi_io.c` on the class website.

It may be very useful during development to be able to print `dist`, `pred` and `known`. The preceding section has a brief discussion of how one might collect the distributed arrays onto process 0 for output.

Extra Help

I have office hours Mondays and Fridays from 4:45 to 5:45 and Wednesdays from 10:30 to 11:30. Erika has office hours on Thursdays from 3 to 4, and Mark has office hours on Thursdays from 4:30 to 5:30.

Submission

You should upload your program to the `p3` subdirectory of your SVN repository. After you first create your source file, type

```
$ svn add mpi_dijk.c
$ svn commit mpi_dijk.c -m "create program 3 source file"
```

Be sure to commit your final version of the source code by 6 pm on Wednesday, October 28.

```
$ svn commit mpi_dijk.c -m "final testing completed"
```

Grading

1. Correctness will be 75% of your grade. Does your program find the correct shortest paths given correct input? Does it correctly store the various distributed data structures?

2. Documentation will be 10% of your grade. Does your header documentation include the author's name, the purpose of the program, and a description of how to use the program? Are the identifiers meaningful? Are any obscure constructs clearly explained? Does the function header documentation explain the purpose of the function, its arguments and its return value?
3. Source format will be 5% of your grade. Is the indentation consistent? Have blank lines been used so that the program is easy to read? Is the use of capitalization consistent? Are there any lines of source that are longer than 80 characters (i.e., wrap around the screen)?
4. Quality of solution will be 10% of your grade. Are any of your functions more than 40 lines (not including blank lines, curly braces, or comments)? Are there multipurpose functions? Is your solution too clever — i.e., has the solution been condensed to the point where it's incomprehensible? Do your functions do unnecessary computations?

Academic Honesty

Remember that you can discuss the program with your classmates, but you cannot look at anyone else's pseudo-code or source code. (Of course, this includes code that you might find on the internet.)